

Lab 3: More Linear Regression in R

Yifan Jin

9/23/2020

Today's Objectives

- Fitting multiple regression and fitting regression with transformed predictors (higher order pattern of the predictors)
- Including categorical variables in linear regression
- Brief introduction to interactions

```
library(ISLR)
```

```
data(Auto)
```

Fitting multiple regression with `lm()`

Suppose we also think vehicle weight will influence fuel efficiency. To fit a multiple linear regression, simply add more variables to the formula in `lm()`.

```
attach(Auto)
```

```
# fit a multiple linear regression  
mlm.fit <- lm(mpg~horsepower+weight)
```

```
# view summary output  
summary(mlm.fit)
```

```
##  
## Call:  
## lm(formula = mpg ~ horsepower + weight)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -11.0762  -2.7340  -0.3312   2.1752  16.2601   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  45.6402108  0.7931958  57.540  < 2e-16 ***  
## horsepower  -0.0473029  0.0110851  -4.267  2.49e-05 ***  
## weight      -0.0057942  0.0005023 -11.535  < 2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 4.24 on 389 degrees of freedom  
## Multiple R-squared:  0.7064, Adjusted R-squared:  0.7049   
## F-statistic: 467.9 on 2 and 389 DF,  p-value: < 2.2e-16
```

A few things to note:

- Both predictors appear to have an influence on fuel efficiency
- The R-squared value has increased after adding another predictor

- The estimate of the coefficient for horsepower has changed

Suppose we want to add even more parameters to the model. The formula shortcut to add all variables is ‘.’ (Warning: the ‘.’ command is not compatible with ‘attach’, you’ll have to specify the dataset in the lm call)

However in this case, the data contains non-numeric variables ‘origin’ and ‘name’ which we can remove from the model by subtracting them in the formula.

```
# fit a multiple linear regression
mlm.fit.2 <- lm(mpg~.-origin-name,data=Auto)

# view summary output
summary(mlm.fit.2)

##
## Call:
## lm(formula = mpg ~ . - origin - name, data = Auto)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.6927 -2.3864 -0.0801  2.0291 14.3607
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.454e+01  4.764e+00  -3.051  0.00244 **
## cylinders    -3.299e-01  3.321e-01  -0.993  0.32122
## displacement  7.678e-03  7.358e-03   1.044  0.29733
## horsepower   -3.914e-04  1.384e-02  -0.028  0.97745
## weight       -6.795e-03  6.700e-04 -10.141 < 2e-16 ***
## acceleration  8.527e-02  1.020e-01   0.836  0.40383
## year          7.534e-01  5.262e-02  14.318 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.435 on 385 degrees of freedom
## Multiple R-squared:  0.8093, Adjusted R-squared:  0.8063
## F-statistic: 272.2 on 6 and 385 DF,  p-value: < 2.2e-16
```

Fitting regression with transformed predictors with lm()

Recall the plot we produced of residuals vs fitted values in our simple linear regression (y=mpg, x=horsepower). We can see a clear quadratic pattern in this plot which implies we could improve the fit by adding a quadratic term to our model.

There are three ways to do this:

- (1) Create a new column in the data frame

```
# create a column containing the squared horsepower
Auto$horsepower2 <- horsepower^2

# fit a multiple linear regression on horsepower and horsepower^2
qm.fit <- lm(mpg~horsepower + horsepower2,data=Auto)

# view basic output
qm.fit
```

```
##
## Call:
## lm(formula = mpg ~ horsepower + horsepower2, data = Auto)
##
## Coefficients:
## (Intercept)    horsepower    horsepower2
##    56.900100    -0.466190     0.001231
```

(2) Do the transformation directly in `lm()`

```
# fit a multiple linear regression
qm.fit2 <- lm(mpg~horsepower+I(horsepower^2))

# view basic output
qm.fit2
```

```
##
## Call:
## lm(formula = mpg ~ horsepower + I(horsepower^2))
##
## Coefficients:
##      (Intercept)      horsepower  I(horsepower^2)
##      56.900100      -0.466190       0.001231
```

(3) Use the `poly()` command to fit a higher order polynomial model


```
# fit a multiple linear regression
qm.fit3 <- lm(mpg~poly(horsepower,2))

# view basic output
qm.fit3
```

```
##
## Call:
## lm(formula = mpg ~ poly(horsepower, 2))
##
## Coefficients:
##      (Intercept)  poly(horsepower, 2)1  poly(horsepower, 2)2
##      23.45      -120.14       44.09
```

```
# the poly command uses orthonormal polynomials, so it produces
# different coefficients but the same fitted values
plot(fitted.values(qm.fit),fitted.values(qm.fit3),
     main="Comparison of fitted values")
abline(a=0,b=1,lty=2)
```

Comparison of fitted values




The `poly` command can be used to fit even higher order polynomial models.

After adding a quadratic term, we have removed the pattern in the residual vs fitted value plot.

```
# quadratic model residuals v fitted values plot
y.hat <- predict(qm.fit)
e.hat <- residuals(qm.fit)
plot(y.hat,e.hat,
     xlab="Fitted values",ylab="Residuals",
     main="Residuals v Fitted values (quadratic model)")
abline(h=0,lty=3)
```

Residuals v Fitted values (quadratic model)



Categorical predictors

2-level categorical predictors

We first consider a binary categorical predictor into simple linear regression.

- A typical appearance for such predictor is the answer from ‘Yes’/‘No’ questionnaire.
- For example, let ‘ x ’ be a binary predictor indicating whether you are from Michigan or not.
- The convention is to treat them as 0/1 indicators. That is, we label ‘yes’ as 1, while ‘no’ as 0. (Sometimes the other way around..)

Consider the following model

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i,$$

where x_i only takes 0 or 1, indicating the binary predictor. This means

$$y_i = \begin{cases} \beta_0 + \beta_1 + \varepsilon_i, & \text{if } x_i = 1 \\ \beta_0 + \varepsilon_i, & \text{if } x_i = 0 \end{cases}$$

- What’s the interpretation of β_0 and β_1 ?
- β_1 : The average difference in Y between the two categories.
- β_0 : The average of Y for the baseline group (Here baseline is $x_i = 0$).

Following the same strategy in lecture, we can derive the expression for the estimator $\hat{\beta}_0$ and $\hat{\beta}_1$. Let \bar{Y}_1 be the sample average of y_i among those $x_i = 1$, while \bar{Y}_0 be the sample average of y_i among $x_i = 0$.

$$\begin{aligned} \hat{\beta}_0 &= \bar{Y}_0 \\ \hat{\beta}_1 &= \bar{Y}_1 - \bar{Y}_0 \end{aligned}$$

- Does those expressions seems intuitive?
- What would happen if we label **Yes** as 2 and **No** as 0?


To see a real-world application, we turn to the **Credit** dataset from the **ISLR** package. The data set records information on credit card clients. We first try to model **Balance** using a binary predictor **Student**.

- **Balance** records the average credit balance for a client
- **Student** is binary, representing whether the user is a student or not.
- A summary of the variable **Student** shows it's indeed categorical.

```
library(ISLR)
data(Credit)
summary(Credit$Student)
```

```
## No Yes
## 360 40
```

```
plot(Balance~Student, data = Credit)
```



Then we can fit a simple linear regression model of **Balance~Student**.

```
fit1 = lm(Balance~Student, data = Credit)
summary(fit1)
```

```
##
## Call:
## lm(formula = Balance ~ Student, data = Credit)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -876.82 -458.82  -40.87   341.88 1518.63
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    480.37      23.43   20.50  < 2e-16 ***
```

```
## StudentYes      396.46      74.10      5.35 1.49e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 444.6 on 398 degrees of freedom
## Multiple R-squared:  0.06709,    Adjusted R-squared:  0.06475
## F-statistic: 28.62 on 1 and 398 DF,  p-value: 1.488e-07
```

- Is the coefficient for ‘Student’ significant?
- While R implement the regression above, how do we know if ‘Yes’ = 1 or ‘No’ = 1?

We can easily verify that the estimated values have the desired interpretation.

```
beta0.hat = mean(Credit$Balance[Credit$Student == 'No'])
beta1.hat = mean(Credit$Balance[Credit$Student == 'Yes']) - mean(Credit$Balance[Credit$Student == 'No'])
c(beta0.hat,beta1.hat)
```

```
## [1] 480.3694 396.4556
```

```
coef(fit1)
```

```
## (Intercept) StudentYes
##      480.3694      396.4556
```

If we would like to put ‘Yes’ as our baseline instead, the function `relevel` comes handy.

```
fit2 = lm(Balance~relevel(Student, ref = 'Yes'), data = Credit)
summary(fit2)
```

```
##
## Call:
## lm(formula = Balance ~ relevel(Student, ref = "Yes"), data = Credit)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -876.83 -458.83  -40.87   341.88 1518.63
##
## Coefficients:
##                                Estimate Std. Error t value Pr(>|t|)
## (Intercept)                   876.8      70.3    12.47 < 2e-16 ***
## relevel(Student, ref = "Yes")No -396.5      74.1    -5.35 1.49e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 444.6 on 398 degrees of freedom
## Multiple R-squared:  0.06709,    Adjusted R-squared:  0.06475
## F-statistic: 28.62 on 1 and 398 DF,  p-value: 1.488e-07
```

Another way of doing this is adjusting the order of the levels first.

```
levels(Credit$Student)
```

```
## [1] "No"  "Yes"
```

```
Credit$MyStudent = factor(Credit$Student, levels = c("Yes", "No"))
levels(Credit$MyStudent)
```

```
## [1] "Yes" "No"
```

```
fit2.1 = lm(Balance~MyStudent , data = Credit)
summary(fit2.1)
```

```
##
## Call:
## lm(formula = Balance ~ MyStudent, data = Credit)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -876.83 -458.83  -40.87   341.88 1518.63
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    876.8       70.3   12.47 < 2e-16 ***
## MyStudentNo   -396.5       74.1    -5.35 1.49e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 444.6 on 398 degrees of freedom
## Multiple R-squared:  0.06709,    Adjusted R-squared:  0.06475
## F-statistic: 28.62 on 1 and 398 DF,  p-value: 1.488e-07
```

- The `level` function reveals all the possible levels of a categorical variable **in order**.
- The `factor` function creates a categorical variable in R with the specified level names **in order**.

Multiple regression

Now let's try to include some other continuous variable in the model.

```
fit3 = lm(Balance~Student + Income + Rating + Age, data = Credit)
summary(fit3)
```

```
##
## Call:
## lm(formula = Balance ~ Student + Income + Rating + Age, data = Credit)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -217.606  -79.887   -8.163    62.680   292.009
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -547.30470    21.46064  -25.503 <2e-16 ***
## StudentYes   417.50564    17.17164   24.314 <2e-16 ***
## Income       -7.79773     0.24218  -32.198 <2e-16 ***
## Rating        3.98073     0.05458   72.927 <2e-16 ***
## Age         -0.62418     0.30407   -2.053  0.0408 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 102.9 on 395 degrees of freedom
## Multiple R-squared:  0.9504, Adjusted R-squared:  0.9499
## F-statistic: 1892 on 4 and 395 DF,  p-value: < 2.2e-16
```

- How would you interpret the coefficient for `Student` now?


Multi-level categorical predictor

Now let's consider the case with more than two levels. First we take a look at the variable `Ethnicity`.

```
summary(Credit$Ethnicity)
```

```
## African American      Asian      Caucasian  
##                99                102                199
```

```
plot(Balance ~ Ethnicity , data = Credit)
```



- There are three levels.
- There does not seem to be a significant relationship between the balance and ethnicity.
- Would indexing `Ethnicity` as a integer of 0/1/2 work like the binary case?

To incorporate this, we introduce a set of ‘dummy’ variables. Suppose x has d distinct levels denoted as f_1, \dots, f_d (They are not numeric!). Let z_1, \dots, z_{d-1} be $d - 1$ binary (dummy) variables that are:

$$z_k = \begin{cases} 1, & \text{if } x = f_k \\ 0, & \text{Otherwise} \end{cases}$$

Then we fit the regression as if we use those $d - 1$ dummy variables z_1, \dots, z_{d-1} instead of the original categorical variable x in the linear regression.

- The dummy z_k indicates whether x takes the k -th value f_k or not.
- Why are there only $d - 1$ dummy variables?
-
- At most one of the dummies z_k will equal to 1.
- When all the $d - 1$ dummies are 0, we immediately know $x = f_d$.
- The level f_d serves as the baseline.

Now let's see how that works in practice.

```
fit4 = lm(Balance~ Ethnicity, data = Credit)
summary(fit4)

##
## Call:
## lm(formula = Balance ~ Ethnicity, data = Credit)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -531.00 -457.08  -63.25   339.25 1480.50
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      531.00      46.32  11.464 <2e-16 ***
## EthnicityAsian    -18.69      65.02  -0.287   0.774
## EthnicityCaucasian -12.50      56.68  -0.221   0.826
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 460.9 on 397 degrees of freedom
## Multiple R-squared:  0.0002188, Adjusted R-squared:  -0.004818
## F-statistic: 0.04344 on 2 and 397 DF,  p-value: 0.9575
```

- Which level is the baseline?
- Now, the coefficients represents the average difference with the **baseline** level.
- Does it appear to be a significant difference between each ethnicity?

Alternatively, we can create our own 0/1 valued dummy variables.

```
Credit$AfAmer <- as.numeric(Credit$Ethnicity == 'African American')
Credit$Asian <- as.numeric(Credit$Ethnicity == 'Asian')
Credit$Cauc <- as.numeric(Credit$Ethnicity == 'Caucasian')
table(Credit$AfAmer)
```

```
##
##      0      1
## 301  99
```

- The == operator in R compares the left and right hand side, returning a logical value TRUE/FALSE.
- In our case, it returns a vector of the same length as `Credit$Ethnicity`.
- The `as.numeric` convert those logical variable to 1/0.

Suppose we use **African American** as the baseline level. Fitting a linear regression using our dummies gives the same result.

```
fit5 = lm(Balance~ Asian + Cauc, data = Credit)
coef(fit5)
```

```
## (Intercept)      Asian      Cauc
##   531.00000   -18.68627   -12.50251
```

```
coef(fit4)
```

```
##      (Intercept)      EthnicityAsian EthnicityCaucasian
##      531.00000      -18.68627      -12.50251
```

Of course we can involve more predictors in the model.

```
fit6 = lm(Balance~ Ethnicity + Student + Income + Rating + Age, data = Credit)
summary(fit6)
```

```
##
## Call:
## lm(formula = Balance ~ Ethnicity + Student + Income + Rating +
##     Age, data = Credit)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -218.14  -82.41  -11.24   64.17  291.85
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -558.91663    23.82767  -23.457  <2e-16 ***
## EthnicityAsian     19.11395    14.58032   1.311   0.1906
## EthnicityCaucasian   9.24769    12.68082   0.729   0.4663
## StudentYes        416.82254    17.20402  24.228  <2e-16 ***
## Income           -7.80138     0.24238 -32.186  <2e-16 ***
## Rating             3.98304     0.05465  72.889  <2e-16 ***
## Age              -0.59631     0.30492  -1.956   0.0512 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 102.9 on 393 degrees of freedom
## Multiple R-squared:  0.9506, Adjusted R-squared:  0.9499
## F-statistic: 1261 on 6 and 393 DF,  p-value: < 2.2e-16
```

- Note the interpretation of the ‘intercept’ terms now involves the baseline of Ethnicity and Student.

Interactions

We will briefly introduce the interaction effect. An interaction term between two variables x_1 and x_2 means the product term $x_1 \cdot x_2$. For example, a linear regression with predictor x_1 , x_2 and their interaction is:

$$y = \beta_0 + x_1\beta_1 + x_2\beta_2 + x_1x_2\beta_3 + \varepsilon$$

- This is particularly interesting when x_1 is categorical (binary).
- Intuitively, the marginal effect of x_2 depend on the level of x_1 .

Let’s look at the interaction between Student and Rating.

```
fit7 = lm(Balance ~ Income + Student + Rating + Student:Rating , data = Credit)
summary(fit7)
```

```
##
## Call:
## lm(formula = Balance ~ Income + Student + Rating + Student:Rating,
##     data = Credit)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -217.930  -78.225   -5.887   66.799  284.319
##
## Coefficients:
```

```
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -568.38426   14.07027 -40.396 < 2e-16 ***
## Income        -7.91290    0.23684 -33.410 < 2e-16 ***
## StudentYes    271.73432   43.97212   6.180 1.6e-09 ***
## Rating         3.95653    0.05456  72.518 < 2e-16 ***
## StudentYes:Rating 0.41548    0.11463   3.624 0.000327 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 101.8 on 395 degrees of freedom
## Multiple R-squared:  0.9515, Adjusted R-squared:  0.951
## F-statistic: 1937 on 4 and 395 DF, p-value: < 2.2e-16
```

- The colon : represents the two-way interaction
- Alternatively we can use `Student*Rating` to represent both the main effect and the interaction: `Student + Rating + Student:Rating`.

```
fit8 = lm(Balance ~ Income + Student*Rating , data = Credit)
summary(fit8)
```

```
##
## Call:
## lm(formula = Balance ~ Income + Student * Rating, data = Credit)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -217.930  -78.225   -5.887    66.799   284.319
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -568.38426   14.07027 -40.396 < 2e-16 ***
## Income        -7.91290    0.23684 -33.410 < 2e-16 ***
## StudentYes    271.73432   43.97212   6.180 1.6e-09 ***
## Rating         3.95653    0.05456  72.518 < 2e-16 ***
## StudentYes:Rating 0.41548    0.11463   3.624 0.000327 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 101.8 on 395 degrees of freedom
## Multiple R-squared:  0.9515, Adjusted R-squared:  0.951
## F-statistic: 1937 on 4 and 395 DF, p-value: < 2.2e-16
```